

▼ CS156 (Introduction to AI), Spring 2022

Homework 8 submission

Roster Name: Bernard Tan

Preferred Name (if different): Bernard

Student ID: 015215317

Email address: bernard.tan@sjsu.edu

Any special notes or anything you would like to communicate to me about this homework submission goes in here.

▼ References and sources

List all your references and sources here. This includes all sites/discussion boards/blogs/posts/etc. where you grabbed some code examples.

▼ Solution

▼ Load libraries and set random number generator seed

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import os
import matplotlib.pyplot as plt
from skimage import io
import numpy as np
```

```
np.random.seed(42)
```

```
image_size = (180, 180)
batch_size = 32
```

```
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "./flowers/training",
    validation_split=0.2,
    subset="training",
    seed=42,

    labels='inferred',
    label_mode='categorical',

    image_size=image_size,
    batch_size=batch_size,
)
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "./flowers/training",
    validation_split=0.2,
    subset="validation",
    seed=42,

    labels='inferred',
    label_mode='categorical',

    image_size=image_size,
    batch_size=batch_size,
)
test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "./flowers/test",
    seed=42,

    labels='inferred',
    label_mode='categorical',

    image_size=image_size,
    batch_size=1,
)
```

```
Found 3456 files belonging to 5 classes.
Using 2765 files for training.
Found 3456 files belonging to 5 classes.
Using 691 files for validation.
Found 861 files belonging to 5 classes.
```

```
data_augmentation = keras.Sequential(
    [
        layers.experimental.preprocessing.RandomFlip("horizontal"),
        layers.experimental.preprocessing.RandomRotation(0.1),
    ]
)
```

```
plt.figure(figsize=(10, 10))
```

```

for images, _ in train_ds.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")

```



```

train_ds = train_ds.prefetch(buffer_size=32)
val_ds = val_ds.prefetch(buffer_size=32)

def make_model(input_shape, num_classes):
    inputs = keras.Input(shape=input_shape)
    # Image augmentation block
    x = data_augmentation(inputs)

    # Entry block
    x = layers.experimental.preprocessing.Rescaling(1.0 / 255)(x)
    x = layers.Conv2D(32, 3, strides=2, padding="same")(x)
    x = layers.BatchNormalization()(x)

```

```

x = layers.Activation("relu")(x)

x = layers.Conv2D(64, 3, padding="same")(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)

previous_block_activation = x # Set aside residual

for size in [128, 256, 512, 728]:
    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same")(x)
    x = layers.BatchNormalization()(x)

    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same")(x)
    x = layers.BatchNormalization()(x)

    x = layers.MaxPooling2D(3, strides=2, padding="same")(x)

    # Project residual
    residual = layers.Conv2D(size, 1, strides=2, padding="same")(
        previous_block_activation
    )
    x = layers.add([x, residual]) # Add back residual
    previous_block_activation = x # Set aside next residual

x = layers.SeparableConv2D(1024, 3, padding="same")(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)

x = layers.GlobalAveragePooling2D()(x)
if num_classes == 2:
    activation = "sigmoid"
    units = 1
else:
    activation = "softmax"
    units = num_classes

x = layers.Dropout(0.5)(x)
outputs = layers.Dense(units, activation=activation)(x)
return keras.Model(inputs, outputs)

```

```

model = make_model(input_shape=image_size + (3,), num_classes=5)
#keras.utils.plot_model(model, show_shapes=True)
model.summary()

```

add_1 (Add)	(None, 23, 23, 256)	0	max_pooling2d_1[0][0] conv2d_3[0][0]
activation_6 (Activation)	(None, 23, 23, 256)	0	add_1[0][0]

separable_conv2d_4 (SeparableCo	(None, 23, 23, 512)	133888	activation_6[0][0]
batch_normalization_6 (BatchNor	(None, 23, 23, 512)	2048	separable_conv2d_4[0]
activation_7 (Activation)	(None, 23, 23, 512)	0	batch_normalization_
separable_conv2d_5 (SeparableCo	(None, 23, 23, 512)	267264	activation_7[0][0]
batch_normalization_7 (BatchNor	(None, 23, 23, 512)	2048	separable_conv2d_5[0]
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 512)	0	batch_normalization_
conv2d_4 (Conv2D)	(None, 12, 12, 512)	131584	add_1[0][0]
add_2 (Add)	(None, 12, 12, 512)	0	max_pooling2d_2[0][0] conv2d_4[0][0]
activation_8 (Activation)	(None, 12, 12, 512)	0	add_2[0][0]
separable_conv2d_6 (SeparableCo	(None, 12, 12, 728)	378072	activation_8[0][0]
batch_normalization_8 (BatchNor	(None, 12, 12, 728)	2912	separable_conv2d_6[0]
activation_9 (Activation)	(None, 12, 12, 728)	0	batch_normalization_
separable_conv2d_7 (SeparableCo	(None, 12, 12, 728)	537264	activation_9[0][0]
batch_normalization_9 (BatchNor	(None, 12, 12, 728)	2912	separable_conv2d_7[0]
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 728)	0	batch_normalization_
conv2d_5 (Conv2D)	(None, 6, 6, 728)	373464	add_2[0][0]
add_3 (Add)	(None, 6, 6, 728)	0	max_pooling2d_3[0][0] conv2d_5[0][0]
separable_conv2d_8 (SeparableCo	(None, 6, 6, 1024)	753048	add_3[0][0]
batch_normalization_10 (BatchNo	(None, 6, 6, 1024)	4096	separable_conv2d_8[0]
activation_10 (Activation)	(None, 6, 6, 1024)	0	batch_normalization_
global_average_pooling2d (Globa	(None, 1024)	0	activation_10[0][0]
dropout (Dropout)	(None, 1024)	0	global_average_pooli
dense (Dense)	(None, 5)	5125	dropout[0][0]
=====			
Total params: 2,786,749			
Trainable params: 2,778,013			
Non-trainable params: 8,736			

epochs = 20

```

callbacks = [
    keras.callbacks.ModelCheckpoint("save_at_{epoch}.h5"),
]
model.compile(
    optimizer=keras.optimizers.Adam(1e-3),
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)
model.fit(
    train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds,
)

Epoch 1/20
87/87 [=====] - 262s 3s/step - loss: 1.2626 - accuracy: 0.5342
Epoch 2/20
/Users/nicholas_dave/opt/anaconda3/lib/python3.8/site-packages/keras/utils/generic_utils.py:30: UserWarning: Custom mask layers require a config and must override the get_config method.
warnings.warn('Custom mask layers require a config and must override the get_config method.')
87/87 [=====] - 268s 3s/step - loss: 1.0034 - accuracy: 0.6253
Epoch 3/20
87/87 [=====] - 262s 3s/step - loss: 0.8878 - accuracy: 0.6723
Epoch 4/20
87/87 [=====] - 263s 3s/step - loss: 0.8245 - accuracy: 0.6890
Epoch 5/20
87/87 [=====] - 262s 3s/step - loss: 0.7971 - accuracy: 0.7016
Epoch 6/20
87/87 [=====] - 264s 3s/step - loss: 0.7587 - accuracy: 0.7226
Epoch 7/20
87/87 [=====] - 254s 3s/step - loss: 0.6805 - accuracy: 0.7392
Epoch 8/20
87/87 [=====] - 257s 3s/step - loss: 0.6724 - accuracy: 0.7552
Epoch 9/20
87/87 [=====] - 254s 3s/step - loss: 0.6049 - accuracy: 0.7790
Epoch 10/20
87/87 [=====] - 256s 3s/step - loss: 0.5778 - accuracy: 0.7892
Epoch 11/20
87/87 [=====] - 259s 3s/step - loss: 0.5569 - accuracy: 0.7920
Epoch 12/20
87/87 [=====] - 256s 3s/step - loss: 0.5252 - accuracy: 0.8004
Epoch 13/20
87/87 [=====] - 240s 3s/step - loss: 0.5277 - accuracy: 0.8098
Epoch 14/20
87/87 [=====] - 257s 3s/step - loss: 0.4938 - accuracy: 0.8145
Epoch 15/20
87/87 [=====] - 256s 3s/step - loss: 0.4899 - accuracy: 0.8137
Epoch 16/20
87/87 [=====] - 257s 3s/step - loss: 0.4733 - accuracy: 0.8181
Epoch 17/20
87/87 [=====] - 255s 3s/step - loss: 0.4302 - accuracy: 0.8401
Epoch 18/20
87/87 [=====] - 255s 3s/step - loss: 0.4362 - accuracy: 0.8445
Epoch 19/20
87/87 [=====] - 255s 3s/step - loss: 0.4000 - accuracy: 0.8499
Epoch 20/20

```

87/87 [=====] - 256s 3s/step - loss: 0.3901 - accuracy: 0.8590
 <keras.callbacks.History at 0x7ff4da34d8b0>

```

true_labels = []
predicted_labels = []
#x = image, y = label
for x, y in test_ds:
    pred = model.predict(x)
    true_labels.append(np.where(y == 1.)[1][0])
    predicted_labels.append(np.where(pred == np.amax(pred))[1][0])

def plot_confusion_matrix(cm,
                          target_names,
                          title='Confusion matrix',
                          cmap=None,
                          normalize=True):

    import matplotlib.pyplot as plt
    import numpy as np
    import itertools

    accuracy = np.trace(cm) / float(np.sum(cm))
    misclass = 1 - accuracy

    if cmap is None:
        cmap = plt.get_cmap('Blues')

    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()

    if target_names is not None:
        tick_marks = np.arange(len(target_names))
        plt.xticks(tick_marks, target_names, rotation=45)
        plt.yticks(tick_marks, target_names)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]


    thresh = cm.max() / 1.5 if normalize else cm.max() / 2
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        if normalize:
            plt.text(j, i, "{:0.4f}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")
        else:
            plt.text(j, i, "{:,}".format(cm[i, j]),

```

```
horizontalalignment="center",
color="white" if cm[i, j] > thresh else "black")
```

```
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label\naccuracy={:0.4f}; misclass={:0.4f}'.format(accuracy, misclas
plt.show())
```

```
plot_confusion_matrix(cm = tf.math.confusion_matrix(labels=true_labels, predictions=predicted
normalize      = False,
target_names   = ['Daisy', 'Dandelion', 'Rose', 'Sunflower', 'Tulip'],
title          = "Flower Prediction on Test Set")
```



```
def target_translator (input_number) :
```

```
    if (input_number == 0) :
        return 'Daisy'
```

```
    elif (input_number == 1) :
        return 'Dandelion'
```

```
    elif (input_number == 2) :
        return 'Rose'
```



```
elif (input_number == 3) :  
    return 'Sunflower'  
  
elif (input_number == 4) :  
    return 'Tulip'
```

```
breaker = 0
```

```
for counter in range (100) :  
    if (true_labels[counter] != predicted_labels[counter]) :  
  
        plt.title(target_translator(true_labels[counter]) + ' is predicted as ' + target_trar  
  
        plt.imshow(images[counter].numpy().astype("uint8"))  
  
        plt.show()  
  
        breaker = breaker + 1  
  
if (breaker > 2) :  
    break
```

